


GENERATIVE ARTIFICIAL INTELLIGENCE

AIML303

Practical file



AMITY SCHOOL OF ENGINEERING AND TECHNOLOGY

AMITY UNIVERSITY, UTTAR PRADESH

Submitted To:

Dr. Ritesh Maurya

Submitted By: Farhan Khan

A2305221537

INDEX

S.no.	Program	Date	Sign
1.	Implement perceptron using Pytorch		
2.	Implement MLP using Pytorch		
3.	Implement MLP for XOR gate using Pytorch		
4.	Implement simple NN for MNIST digit classification using Pytorch.		
5.	Implement simple NN for MNIST digit classification using Keras and TF		
6.	Implement CNN using Keras for CIFAR10 classification		
7.	Transfer Learning for sugarcane disease classification		
8.	Image denoising using autoencoder		
9.	Image generation using convolutional GAN		
10.	Simple audio recognition		
11.	Text generation using RNN		

EXPERIMENT - 1

Aim:

Implement perceptron using pytorch.

Experiment Conducted:

```
import numpy as np

[ ] #define perceptron parameters
input_size = 2
weights = np.random.rand(input_size)
bias = np.random.rand(1)
# learning_rate = 0.01

[ ] #define the step function
def step_function(x): return 1 if x > 0 else 0

[ ] # define the perceptron function
def perceptron(inputs,weights,bias):
    linear_combination = np.dot(inputs, weights) + bias
    output = step_function(linear_combination)
    return output

[ ] #generate random inputs
inputs = np.random.rand(input_size)

[ ] #apply the perceptron
output = perceptron(inputs,weights,bias)

[ ] #display results
print("Inputs:",inputs)
print("Weights:",weights)
print("Bias:",bias)
print("output",output)
```

Output:

```
Inputs: [0.90863384 0.6993004 ]
Weights: [0.18947875 0.87561118]
Bias: [0.03618943]
output 1
```

Conclusion: Hence, the experiment is conducted successfully and the results have been verified.

EXPERIMENT - 2

Aim:

Implement MLP using Pytorch.

Experiment Conducted:

```
import torch
import torch.nn as nn
import torch.optim as optim

[ ] #Define the neural network architecture
class MLP(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super(MLP, self).__init__()
        self.fc1 = nn.Linear(input_size,hidden_size)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(hidden_size, output_size)

    def forward(self,x):
        x=self.fc1(x)
        x=self.relu(x)
        x=self.fc2(x)
        return x

[ ] # set random seed for reproducibility
torch.manual_seed(42)

<torch._C.Generator at 0x7959681fdb0>

[ ] # Instantiate the model
input_size = 10
hidden_size = 20
output_size = 1
model = MLP(input_size,hidden_size,output_size)

[ ] # Define the loss function and optimizer
criterion = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=0.5)

[ ] #dummy input data and arget
input_data=torch.randn(100,input_size)
print("input_data_shape",input_data.shape)
target=torch.randn(100,output_size)
#print(target)

input_data_shape torch.Size([100, 10])
```

```
[ ] #Training loop
num_epochs=4000
for epoch in range(num_epochs):
    #Forward pass
    output= model(input_data)
    loss= criterion(output,target)

[ ] #Backward pass and optimization
optimizer.zero_grad()
#Zero the gradients to avoid accumulation
loss.backward()
#Backpropagation
optimizer.step()
#Update weights
#Print the loss every 100 epochs
if (epoch +1)%100==0:
    print(f'Epoch[(epoch +1)/(num_epochs)],Loss:{loss.item():.4f}')

Epoch[4000/4000], Loss:1.1767

[ ] #Test the trained model
test_input=torch.randn(1,input_size)
print("test_input",test_input.shape)
with torch.no_grad():
    predicted_output=model(test_input)
    print(f'Test Input:{test_input.numpy()}')
    print(f'Predicted Output:{predicted_output.numpy()}')
```

Output:

```
test_input torch.Size([1, 10])
Test Input:[[ 0.77188647  0.69540155  1.0965257 -0.00797612  1.3109813  0.25582278
  0.24653137 -0.421456 -0.86095834  0.56766146]]
Predicted Output:[[0.2269126]]
```

Conclusion: Hence, the experiment is conducted successfully and the results have been verified.

EXPERIMENT - 3

Aim:

Implement MLP for XOR gate using PyTorch.

Experiment Conducted:

```
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader

[ ] # Define a simple neural network
class SimpleNN(nn.Module):
    def __init__(self):
        super(SimpleNN, self).__init__()
        self.fc1 = nn.Linear(28*28,128)
        self.relu=nn.ReLU()
        self.fc2=nn.Linear(128,10)

    def forward(self,x):
        X=X.view(-1,28*28)
        #print("shape",x.shape)
        X=self.fc1(X)
        X=self.relu(X)
        X=self.fc2(X)
        return X

# Download and load the MNIST dataset
transform=transforms.Compose([transforms.ToTensor(),transforms.Normalize((0.5), (0.5,))])

train_dataset=datasets.MNIST(root='./data',train=True,transform=transform,download=True)
test_dataset=datasets.MNIST(root='./data',train=False,transform=transform,download=True)

# Print the sizes of the train and test datasets
print(f"Train dataset size:{len(train_dataset)}")
print(f"Test dataset size:{len(test_dataset)}")

Train dataset size:60000
Test dataset size:10000

[ ] train_loader = DataLoader(dataset=train_dataset,batch_size=64,shuffle=True)
test_loader = DataLoader(dataset=test_dataset,batch_size=64,shuffle=False)

[ ] # Initialize the model,loss function, and optimizer
model=SimpleNN()
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(),lr=0.001)

# Training loop
num_epochs=10

for epoch in range(num_epochs):
    for i,(images,labels) in enumerate(train_loader):
        optimizer.zero_grad()
        outputs=model(images)
        loss=criterion(outputs,labels)
        loss.backward()
        optimizer.step()

    if (i+1)%100==0:
        print(f'Epoch[{epoch+1}/{num_epochs}],Step[{i+1}/{len(train_loader)}],Loss:{loss.item():.4f}')

Epoch[10/10],Step[800/930],Loss:0.0024
Epoch[10/10],Step[900/930],Loss:0.0109

# Testing the model
model.eval()
correct,total=0,0

with torch.no_grad():
    for images,labels in test_loader:
        outputs=model(images)
        _, predicted = torch.max(outputs.data,1)
        total+=labels.size(0)
        correct+=(predicted==labels).sum().item()

accuracy = correct/total
print(f'Test Accuracy:{accuracy*100:.2f}%')
```

Output:

Test Accuracy:97.30%

Conclusion: Hence, the experiment is conducted successfully and the results have been verified.

EXPERIMENT - 4

Aim:

Implement simple NN for MNIST digit classification using Pytorch.

Experiment Conducted:

```
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
from sklearn.metrics import confusion_matrix, classification_report
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

[ ] # set device(use GPU if available)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Define a simple neural network
class SimpleNN(nn.Module):
    def __init__(self):
        super(SimpleNN, self).__init__()
        self.fc1 = nn.Linear(28*28, 128)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(128, 10)

    def forward(self, x):
        x = x.view(-1, 28*28)
        x = self.fc1(x)
        x = self.relu(x)
        x = self.fc2(x)
        return x

[ ] # Download and load the MNIST dataset
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
])

train_dataset = datasets.MNIST(root='./data', train=True, transform=transform, download=True)
test_dataset = datasets.MNIST(root='./data', train=False, transform=transform, download=True)

train_loader = DataLoader(dataset=train_dataset, batch_size=64, shuffle=True)
test_loader = DataLoader(dataset=test_dataset, batch_size=64, shuffle=False)

[ ] # Initialize the model, loss function, and optimizer
model = SimpleNN().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Training loop
num_epochs = 5

for epoch in range(num_epochs):
    for i, (images, labels) in enumerate(train_loader):
        images, labels = images.to(device), labels.to(device)
        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        if (i+1) % 100 == 0:
            print(f'Epoch [{epoch+1}/{num_epochs}], Step [{i+1}/{len(train_loader)}], Loss: {loss.item():.4f}')

Epoch [5/5], Step [100/938], Loss: 0.6727
Epoch [5/5], Step [500/938], Loss: 0.6853
Epoch [5/5], Step [900/938], Loss: 0.1098

[ ] # Testing the model
model.eval()
y_true, y_pred = [], []

with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        y_true.extend(labels.cpu().numpy())
        y_pred.extend(predicted.cpu().numpy())


[ ] # Convert lists to Numpy arrays
y_true = np.array(y_true)
y_pred = np.array(y_pred)

# Create a confusion matrix
conf_matrix = confusion_matrix(y_true, y_pred)
```

```
# Plot the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", cbar=False,
            xticklabels=np.arange(10), yticklabels=np.arange(10))
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```

```
# Print classification report
print("Classification Report:")
print(classification_report(y_true, y_pred, target_names=[str(i) for i in range(10)]))
```

Output:



Conclusion: Hence, the experiment is conducted successfully and the results have been verified.

EXPERIMENT - 5

Aim:

Implement simple NN for MNIST digit classification using Keras and TF.

Experiment Conducted:

```
# Importing necessary libraries
import numpy as np
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.utils import to_categorical

[ ] # Load and preprocess the MNIST dataset
(x_train, y_train), (x_test, y_test) = mnist.load_data()

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 0s 0us/step

[ ] # Normalize pixel values to be between 0 and 1
x_train = x_train / 255.0
x_test = x_test / 255.0

[ ] # Flatten the 28x28 images into 1D arrays
x_train = x_train.reshape((x_train.shape[0], -1))
x_test = x_test.reshape((x_test.shape[0], -1))

[ ] # One-hot encode the target labels
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)

[ ] # Build the MLP model
model = Sequential()
model.add(Dense(128, input_shape=(784,), activation='relu'))
model.add(Dense(64, activation='relu'))
model.add(Dense(10, activation='softmax'))

[ ] # Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Train the model
history=model.fit(x_train, y_train, epochs=20, batch_size=32, validation_split=0.2)


[ ] # Evaluate the model on the test set
loss, accuracy = model.evaluate(x_test, y_test)
print(f'Test Accuracy: {accuracy * 100:.2f}%')

# Plot training and validation accuracy curves
import matplotlib.pyplot as plt

plt.plot(history.history['accuracy'],label='Training Accuracy')
plt.plot(history.history['val_accuracy'],label='Validation Accuracy')
plt.title('Training and Validation Accuracy Curves')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```

Output:

```
313/313 [=====] - 1s 2ms/step - loss: 0.1188 - accuracy: 0.9778
Test Accuracy: 97.78%
```



Conclusion: Hence, the experiment is conducted successfully and the results have been verified.

EXPERIMENT - 6

Aim:

Implement CNN using Keras for CIFAR10 classification.

Experiment Conducted:

```
import tensorflow as tf
from tensorflow.keras import datasets, layers, models, optimizers
from tensorflow.keras.applications import VGG16
from tensorflow.keras.preprocessing.image import ImageDataGenerator

# Load CIFAR-10 dataset
(train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()

# Normalize pixel values to be between 0 and 1
train_images, test_images = train_images / 255.0, test_images / 255.0

# Define data augmentation
train_datagen = ImageDataGenerator(
    rotation_range=15,
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=True,
)
train_datagen.fit(train_images)

# Load pre-trained VGG16 model without including top dense layers
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(32, 32, 3))

# Freeze all layers in base model
for layer in base_model.layers:
    layer.trainable = False

# Add custom dense layers for CIFAR-10 classification
model = models.Sequential()
model.add(base_model)
model.add(layers.Flatten())
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(10, activation='softmax'))

# Compile the model
model.compile(optimizer=optimizers.Adam(),
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(train_datagen.flow(train_images, train_labels, batch_size=64),
                    steps_per_epoch=len(train_images) / 64, epochs=10,
                    validation_data=(test_images, test_labels))

# Evaluate the model
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f'Test accuracy: {test_acc}')
```

Output:

```
Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 [=====] - 95s 1us/step
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notoop.h5
58889256/58889256 [=====] - 4s 0us/step
Epoch 1/10
781/781 [=====] - 44s 50ms/step - loss: 1.6040 - accuracy: 0.1001 - val_loss: 1.3276 - val_accuracy: 0.1136
Epoch 2/10
781/781 [=====] - 35s 45ms/step - loss: 1.4076 - accuracy: 0.0984 - val_loss: 1.2594 - val_accuracy: 0.1031
Epoch 3/10
781/781 [=====] - 35s 44ms/step - loss: 1.3603 - accuracy: 0.0968 - val_loss: 1.2182 - val_accuracy: 0.1246
Epoch 4/10
781/781 [=====] - 34s 44ms/step - loss: 1.3256 - accuracy: 0.0981 - val_loss: 1.2021 - val_accuracy: 0.0943
Epoch 5/10
781/781 [=====] - 35s 44ms/step - loss: 1.3104 - accuracy: 0.0970 - val_loss: 1.2014 - val_accuracy: 0.1000
Epoch 6/10
781/781 [=====] - 35s 45ms/step - loss: 1.3005 - accuracy: 0.0978 - val_loss: 1.1776 - val_accuracy: 0.0984
Epoch 7/10
781/781 [=====] - 35s 45ms/step - loss: 1.2816 - accuracy: 0.0977 - val_loss: 1.1735 - val_accuracy: 0.1187
Epoch 8/10
781/781 [=====] - 35s 45ms/step - loss: 1.2731 - accuracy: 0.0963 - val_loss: 1.1796 - val_accuracy: 0.1122
Epoch 9/10
781/781 [=====] - 36s 46ms/step - loss: 1.2657 - accuracy: 0.0994 - val_loss: 1.1634 - val_accuracy: 0.0895
Epoch 10/10
781/781 [=====] - 36s 45ms/step - loss: 1.2594 - accuracy: 0.0977 - val_loss: 1.1484 - val_accuracy: 0.1020
313/313 [=====] - 3s 9ms/step - loss: 1.1484 - accuracy: 0.1020
Test accuracy: 0.10199999809265137
```

Conclusion: Hence, the experiment is conducted successfully and the results have been verified.

EXPERIMENT - 7

Aim:

Transfer Learning for sugarcane disease classification.

Experiment Conducted:

```
import os
import torch
import torchvision.transforms as transforms
import torchvision.datasets as datasets
import torchvision.models as models
import torch.optim as optim
import torch.nn as nn
import matplotlib.pyplot as plt

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
])

data_folder = '/kaggle/input/sugarcane-disease-dataset/sugarcane RA/'

dataset = datasets.ImageFolder(root=data_folder, transform=transform)

train_size = int(0.8 * len(dataset))
test_size = len(dataset) - train_size
train_dataset, test_dataset = torch.utils.data.random_split(dataset, [train_size,
test_size])

num_classes = len(dataset.classes)
model = models.vgg16(pretrained=True)
model.fc = nn.Linear(4096, num_classes) # VGG16 has 4096 output features
model = model.to(device)

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)

batch_size = 64
trainloader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size,
shuffle=True, num_workers=4)
testloader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size,
shuffle=False, num_workers=4)

num_epochs = 10
train_acc_history = []
test_acc_history = []
```

```

for epoch in range(num_epochs):
    model.train()
    running_loss = 0.0
    total_train = 0
    correct_train = 0

    for i, (inputs, labels) in enumerate(trainloader):
        inputs, labels = inputs.to(device), labels.to(device)

        optimizer.zero_grad()

        outputs = model(inputs)
        loss = criterion(outputs, labels)

        loss.backward()
        optimizer.step()

        running_loss += loss.item()
        _, predicted_train = torch.max(outputs.data, 1)
        total_train += labels.size(0)
        correct_train += (predicted_train == labels).sum().item()

    train_accuracy = 100 * correct_train / total_train
    train_acc_history.append(train_accuracy)

    model.eval()
    correct_test = 0
    total_test = 0

    with torch.no_grad():
        for inputs, labels in testloader:
            inputs, labels = inputs.to(device), labels.to(device)
            outputs = model(inputs)
            _, predicted_test = torch.max(outputs.data, 1)
            total_test += labels.size(0)
            correct_test += (predicted_test == labels).sum().item()

    test_accuracy = 100 * correct_test / total_test
    test_acc_history.append(test_accuracy)


    print(f'Epoch [{epoch+1}/{num_epochs}], Train Accuracy: {train_accuracy:.2f}%, Test
Accuracy: {test_accuracy:.2f}%')

# Plotting accuracy curves
epochs = range(1, num_epochs + 1)
plt.plot(epochs, train_acc_history, label='Train Accuracy')
plt.plot(epochs, test_acc_history, label='Test Accuracy')
plt.title('Training and Test Accuracy Curves')
plt.xlabel('Epoch')

```

```
plt.ylabel('Accuracy (%)')  
plt.legend()  
plt.show()
```

Output:



Conclusion: Hence, the experiment is conducted successfully and the results have been verified.

EXPERIMENT - 8

Aim:

Image denoising using autoencoder.

Experiment Conducted:

```
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt

[ ] # Load MNIST dataset
(x_train, _), (x_test, _) = tf.keras.datasets.mnist.load_data()

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist-no-11498434/11498434 [=====] - 85 0us/step

[ ] # Normalize and reshape the data
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = np.reshape(x_train, (len(x_train), 28, 28, 1))
x_test = np.reshape(x_test, (len(x_test), 28, 28, 1))

[ ] # Add Gaussian noise to the images
noise_factor = 0.5
x_train_noisy = x_train + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_train.shape)
x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_test.shape)
x_train_noisy = np.clip(x_train_noisy, 0., 1.)
x_test_noisy = np.clip(x_test_noisy, 0., 1.)

[ ] # Define the autoencoder model
input_img = tf.keras.layers.Input(shape=(28, 28, 1))

[ ] # Encoder
x = tf.keras.layers.Conv2D(32, (3, 3), activation='relu', padding='same')(input_img)
x = tf.keras.layers.MaxPooling2D((2, 2), padding='same')(x)
x = tf.keras.layers.Conv2D(32, (3, 3), activation='relu', padding='same')(x)
encoded = tf.keras.layers.MaxPooling2D((2, 2), padding='same')(x)

[ ] # Decoder
x = tf.keras.layers.Conv2D(32, (3, 3), activation='relu', padding='same')(encoded)
x = tf.keras.layers.UpSampling2D((2, 2))(x)
x = tf.keras.layers.Conv2D(32, (3, 3), activation='relu', padding='same')(x)
x = tf.keras.layers.UpSampling2D((2, 2))(x)
decoded = tf.keras.layers.Conv2D(1, (3, 3), activation='sigmoid', padding='same')(x)

autoencoder = tf.keras.models.Model(input_img, decoded)
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

[ ] # Train the autoencoder
autoencoder.fit(x_train_noisy, x_train,
               epochs=50,
               batch_size=128,
               shuffle=True,
               validation_data=(x_test_noisy, x_test))

[ ] # Predict the denoised images
decoded_imgs = autoencoder.predict(x_test_noisy)

313/313 [=====] - 1s 2ms/step

# Plot the original, noisy, and denoised images
n = 10
plt.figure(figsize=(20, 6))
for i in range(n):
    # Original images
    ax = plt.subplot(3, n, i + 1)
    plt.imshow(x_test[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
    if i == n // 2:
        ax.set_title('Original Images')

    # Noisy images
    ax = plt.subplot(3, n, i + 1 + n)
    plt.imshow(x_test_noisy[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
    if i == n // 2:
        ax.set_title('Noisy Input')

    # Denoised images
    ax = plt.subplot(3, n, i + 1 + 2 * n)
    plt.imshow(decoded_imgs[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
    if i == n // 2:
        ax.set_title('Denoised Images')
plt.show()
print(autoencoder.summary())
```

Output:



Conclusion: Hence, the experiment is conducted successfully and the results have been verified.

EXPERIMENT - 9

Aim:

Image generation using convolutional GAN.

Experiment Conducted:

```
[ ] import numpy as np
import matplotlib.pyplot as plt
from keras.models import Sequential, Model
from keras.layers import Input, Dense, Reshape, Flatten, Conv2D, Conv2DTranspose, LeakyReLU
from keras.optimizers import Adam
from keras.datasets import mnist
```

```
[ ] # Load MNIST data
(X_train, _), (X_test, _) = mnist.load_data()

# Normalize data
X_train = (X_train.astype(np.float32) - 127.5) / 127.5
X_train = np.expand_dims(X_train, axis=3)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11490434/11490434 [=====] - 0s 0us/step

```
[ ] # Define generator
def build_generator():
    model = Sequential()
    model.add(Dense(7 * 7 * 128, input_dim=100))
    model.add(LeakyReLU(0.2))
    model.add(Reshape((7, 7, 128)))
    model.add(Conv2DTranspose(64, kernel_size=4, strides=2, padding='same'))
    model.add(LeakyReLU(0.2))
    model.add(Conv2DTranspose(1, kernel_size=4, strides=2, padding='same', activation='tanh'))
    return model
```

```
[ ] # Define discriminator
def build_discriminator():
    model = Sequential()
    model.add(Conv2D(64, kernel_size=4, strides=2, padding='same', input_shape=(28, 28, 1)))
    model.add(LeakyReLU(0.2))
    model.add(Conv2D(128, kernel_size=4, strides=2, padding='same'))
    model.add(LeakyReLU(0.2))
    model.add(Flatten())
    model.add(Dense(1, activation='sigmoid'))
    return model
```

```
[ ] # Compile discriminator
discriminator = build_discriminator()
discriminator.compile(loss='binary_crossentropy', optimizer=Adam(lr=0.0002, beta_1=0.5), metrics=['accuracy'])
```

WARNING:absl: 'lr' is deprecated in Keras optimizer, please use 'learning_rate' or use the legacy optimizer, e.g., tf.keras.optimizers.LegacyAdam.

```
[ ] # Combine generator and discriminator into a single model
generator = build_generator()
z = Input(shape=(100,))
img = generator(z)
discriminator.trainable = False
validity = discriminator(img)
gan = Model(z, validity)
gan.compile(loss='binary_crossentropy', optimizer=Adam(lr=0.0002, beta_1=0.5))
```

WARNING:absl: 'lr' is deprecated in Keras optimizer, please use 'learning_rate' or use the legacy optimizer, e.g., tf.keras.optimizers.LegacyAdam.

```
# Train DCGAN
epochs = 1000
batch_size = 64
for epoch in range(epochs):
    # Select a random batch of images
    idx = np.random.randint(0, X_train.shape[0], batch_size)
    real_images = X_train[idx]
    # Generate fake images
    noise = np.random.normal(0, 1, (batch_size, 100))
    fake_images = generator.predict(noise)
    # Train discriminator
    d_loss_real = discriminator.train_on_batch(real_images, np.ones(batch_size))
    d_loss_fake = discriminator.train_on_batch(fake_images, np.zeros(batch_size))
    d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
    # Train generator
    noise = np.random.normal(0, 1, (batch_size, 100))
    g_loss = gan.train_on_batch(noise, np.ones(batch_size))
    # Print progress
    if epoch % 100 == 0:
        print(f"Epoch: {epoch} \t Discriminator Loss: {d_loss[0]} \t Generator Loss: {g_loss}")
    # Generate images
    noise = np.random.normal(0, 1, (100, 100))
    generated_images = generator.predict(noise)
```

```
2/2 [=====] - 1s 7ms/step
Epoch: 0      Discriminator Loss: 0.6859554648399353      Generator Loss: 0.684131383895874
2/2 [=====] - 0s 5ms/step
2/2 [=====] - 0s 4ms/step
```

```
2/2 [=====] - 0s 4ms/step
2/2 [=====] - 0s 5ms/step
2/2 [=====] - 0s 4ms/step
4/4 [=====] - 0s 30ms/step
```


```
# Display generated images
plt.figure(figsize=(10, 10))
for i in range(100):
    plt.subplot(10, 10, i+1)
    plt.imshow(generated_images[i, :, :, 0], cmap='gray')
    plt.axis("off")
plt.tight_layout()
plt.show()
```

```


# Visualize original images
plt.figure(figsize=(10, 10))
for i in range(100):
    plt.subplot(10, 10, i+1)
    plt.imshow(X_train[i, :, :], cmap='gray')
    plt.axis('off')
plt.tight_layout()
plt.show()

```

Output:



(Generated images)



(Original images)

Conclusion: Hence, the experiment is conducted successfully and the results have been verified.

EXPERIMENT - 10

Aim:

Simple Audio recognition.

Experiment Conducted:

```
[1] pip install -U -q tensorflow tensorflow_datasets

===== 589.8/589.8 MB 1.1 MB/s eta 0:00:00
===== 4.8/4.8 MB 86.7 MB/s eta 0:00:00
===== 2.2/2.2 MB 80.2 MB/s eta 0:00:00
===== 5.5/5.5 MB 74.9 MB/s eta 0:00:00
===== 1.0/1.0 MB 63.8 MB/s eta 0:00:00

ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.
tf-keras 2.15.1 requires tensorflow<2.16,>=2.15, but you have tensorflow 2.16.1 which is incompatible.

[3] DATASET_PATH = 'data/mini_speech_commands'

data_dir = pathlib.Path(DATASET_PATH)
if not data_dir.exists():
    tf.keras.utils.get_file(
        'mini_speech_commands.zip',
        origin='http://storage.googleapis.com/download.tensorflow.org/data/mini_speech_commands.zip',
        extract=True,
        cache_dir='.', cache_subdir='data')

Downloading data from http://storage.googleapis.com/download.tensorflow.org/data/mini_speech_commands.zip
182082353/182082353 ————— 2s 0us/step

[4] commands = np.array(tf.io.gfile.listdir(str(data_dir)))
commands = commands[(commands != 'README.md') & (commands != '.DS_Store')]
print('Commands:', commands)

Commands: ['right' 'yes' 'up' 'left' 'go' 'no' 'stop' 'down']

[5] train_ds, val_ds = tf.keras.utils.audio_dataset_from_directory(
    directory=data_dir,
    batch_size=64,
    validation_split=0.2,
    seed=0,
    output_sequence_length=16000,
    subset='both')

label_names = np.array(train_ds.class_names)
print()
print("label_names:", label_names)

Found 8000 files belonging to 8 classes.
Using 6400 files for training.
Using 1600 files for validation.

label_names: ['down' 'go' 'left' 'no' 'right' 'stop' 'up' 'yes']

train_ds.element_spec

(TensorSpec(shape=(None, 16000), dtype=tf.float32, name=None),
 TensorSpec(shape=(None,), dtype=tf.int32, name=None))

[7] def squeeze(audio, labels):
    audio = tf.squeeze(audio, axis=-1)
    return audio, labels

train_ds = train_ds.map(squeeze, tf.data.AUTOTUNE)
val_ds = val_ds.map(squeeze, tf.data.AUTOTUNE)

[8] test_ds = val_ds.shard(num_shards=2, index=0)
val_ds = val_ds.shard(num_shards=2, index=1)

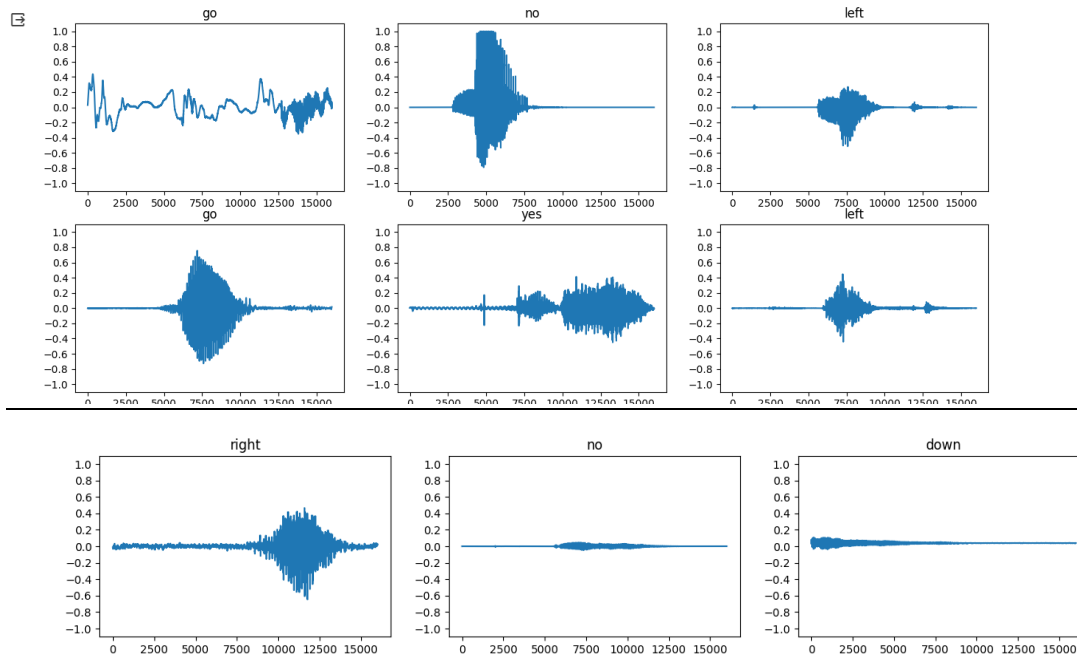
[9] for example_audio, example_labels in train_ds.take(1):
    print(example_audio.shape)
    print(example_labels.shape)

(64, 16000)
(64,)

[10] label_names[[1,1,3,0]]

array(['go', 'go', 'no', 'down'], dtype='<U5')
```

```
plt.figure(figsize=(16, 10))
rows = 3
cols = 3
n = rows * cols
for i in range(n):
    plt.subplot(rows, cols, i+1)
    audio_signal = example_audio[i]
    plt.plot(audio_signal)
    plt.title(label_names[example_labels[i]])
    plt.yticks(np.arange(-1.2, 1.2, 0.2))
    plt.ylim([-1.1, 1.1])
```



```
[12] def get_spectrogram(waveform):
    # Convert the waveform to a spectrogram via a STFT.
    spectrogram = tf.signal.stft(
        waveform, frame_length=255, frame_step=128)
    # Obtain the magnitude of the STFT.
    spectrogram = tf.abs(spectrogram)
    # Add a 'channels' dimension, so that the spectrogram can be used
    # as image-like input data with convolution layers (which expect
    # shape ('batch_size', 'height', 'width', 'channels')).
    spectrogram = spectrogram[..., tf.newaxis]
    return spectrogram
```

```
for i in range(3):
    label = label_names[example_labels[i]]
    waveform = example_audio[i]
    spectrogram = get_spectrogram(waveform)

    print('Label:', label)
    print('Waveform shape:', waveform.shape)
    print('Spectrogram shape:', spectrogram.shape)
    print('Audio playback')
    display.display(display.Audio(waveform, rate=16000))
```

Label: go
Waveform shape: (16000,)
Spectrogram shape: (124, 129, 1)
Audio playback



Label: no
Waveform shape: (16000,)
Spectrogram shape: (124, 129, 1)
Audio playback



Label: left
Waveform shape: (16000,)
Spectrogram shape: (124, 129, 1)
Audio playback



```

[14] def plot_spectrogram(spectrogram, ax):
    if len(spectrogram.shape) > 2:
        assert len(spectrogram.shape) == 3
        spectrogram = np.squeeze(spectrogram, axis=-1)
    # Convert the frequencies to log scale and transpose, so that the time is
    # represented on the x-axis (columns).
    # Add an epsilon to avoid taking a log of zero.
    log_spec = np.log(spectrogram.T + np.finfo(float).eps)
    height = log_spec.shape[0]
    width = log_spec.shape[1]
    X = np.linspace(0, np.size(spectrogram), num=width, dtype=int)
    Y = range(height)
    ax.colormesh(X, Y, log_spec)

```

```


fig, axes = plt.subplots(2, figsize=(12, 8))
timescale = np.arange(waveform.shape[0])
axes[0].plot(timescale, waveform.numpy())
axes[0].set_title('Waveform')
axes[0].set_xlim([0, 16000])

plot_spectrogram(spectrogram.numpy(), axes[1])
axes[1].set_title('Spectrogram')
plt.suptitle(label.title())
plt.show()


```



Left



Spectrogram



```

[16] def make_spec_ds(ds):
    return ds.map(
        map_func=lambda audio, label: (get_spectrogram(audio), label),
        num_parallel_calls=tf.data.AUTOTUNE)

```

```

[17] train_spectrogram_ds = make_spec_ds(train_ds)
    val_spectrogram_ds = make_spec_ds(val_ds)
    test_spectrogram_ds = make_spec_ds(test_ds)

```

```

[18] for example_spectrograms, example_spect_labels in train_spectrogram_ds.take(1):
    break

```

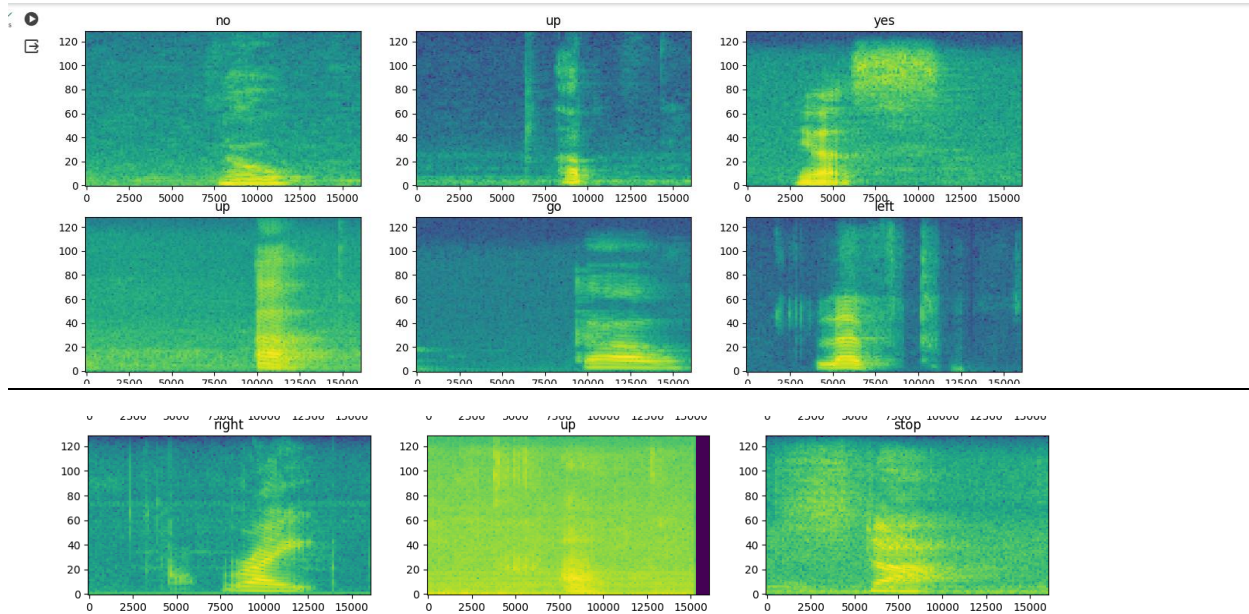
```

rows = 3
cols = 3
n = rows*cols
fig, axes = plt.subplots(rows, cols, figsize=(16, 9))

for i in range(n):
    r = i // cols
    c = i % cols
    ax = axes[r][c]
    plot_spectrogram(example_spectrograms[i].numpy(), ax)
    ax.set_title(label_names[example_spect_labels[i].numpy()])

plt.show()

```



```
[20] train_spectrogram_ds = train_spectrogram_ds.cache().shuffle(10000).prefetch(tf.data.AUTOTUNE)
val_spectrogram_ds = val_spectrogram_ds.cache().prefetch(tf.data.AUTOTUNE)
test_spectrogram_ds = test_spectrogram_ds.cache().prefetch(tf.data.AUTOTUNE)
```

```
input_shape = example_spectrograms.shape[1:]
print('Input shape:', input_shape)
num_labels = len(label_names)

# Instantiate the `tf.keras.layers.Normalization` layer.
norm_layer = layers.Normalization()
# Fit the state of the layer to the spectrograms
# with `Normalization.adapt`.
norm_layer.adapt(data=train_spectrogram_ds.map(map_func=lambda spec, label: spec))

model = models.Sequential([
    layers.Input(shape=input_shape),
    # Downsample the input.
    layers.Resizing(32, 32),
    # Normalize.
    norm_layer,
    layers.Conv2D(32, 3, activation='relu'),
    layers.Conv2D(64, 3, activation='relu'),
    layers.MaxPooling2D(),
    layers.Dropout(0.25),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(num_labels),
])

model.summary()
```

Input shape: (124, 129, 1)
Model: "sequential"

Layer (type)	Output Shape	Param #
resizing (Resizing)	(None, 32, 32, 1)	0
normalization (Normalization)	(None, 32, 32, 1)	3
conv2d (Conv2D)	(None, 30, 30, 32)	320
conv2d_1 (Conv2D)	(None, 28, 28, 64)	18,496
max_pooling2d (MaxPooling2D)	(None, 14, 14, 64)	0
dropout (Dropout)	(None, 14, 14, 64)	0
flatten (Flatten)	(None, 12544)	0
dense (Dense)	(None, 128)	1,605,760
dropout_1 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 8)	1,032

Total params: 1,625,611 (6.20 MB)
Trainable params: 1,625,608 (6.20 MB)
Non-trainable params: 3 (16.00 B)

```
[22] model.compile(
    optimizer=tf.keras.optimizers.Adam(),
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    metrics=['accuracy'],
)

EPOCHS = 10
history = model.fit(
    train_spectrogram_ds,
    validation_data=val_spectrogram_ds,
    epochs=EPOCHS,
    callbacks=tf.keras.callbacks.EarlyStopping(verbose=1, patience=2),
)

Epoch 1/10 ————— 38s 353ms/step - accuracy: 0.2814 - loss: 1.9261 - val_accuracy: 0.6042 - val_loss: 1.3336
Epoch 2/10 ————— 27s 273ms/step - accuracy: 0.5620 - loss: 1.2813 - val_accuracy: 0.7344 - val_loss: 0.9180
Epoch 3/10 ————— 39s 259ms/step - accuracy: 0.6797 - loss: 0.9223 - val_accuracy: 0.7721 - val_loss: 0.7437
Epoch 4/10 ————— 36s 211ms/step - accuracy: 0.7352 - loss: 0.7276 - val_accuracy: 0.7878 - val_loss: 0.6636
Epoch 5/10 ————— 19s 188ms/step - accuracy: 0.7883 - loss: 0.6207 - val_accuracy: 0.8268 - val_loss: 0.5748
Epoch 6/10 ————— 22s 200ms/step - accuracy: 0.8133 - loss: 0.5442 - val_accuracy: 0.8346 - val_loss: 0.5365
Epoch 7/10 ————— 20s 199ms/step - accuracy: 0.8235 - loss: 0.4907 - val_accuracy: 0.8372 - val_loss: 0.5209
Epoch 8/10 ————— 23s 221ms/step - accuracy: 0.8412 - loss: 0.4353 - val_accuracy: 0.8398 - val_loss: 0.4950
Epoch 9/10 ————— 20s 201ms/step - accuracy: 0.8807 - loss: 0.3569 - val_accuracy: 0.8451 - val_loss: 0.4714
Epoch 10/10 ————— 24s 242ms/step - accuracy: 0.8826 - loss: 0.3339 - val_accuracy: 0.8568 - val_loss: 0.4500
```

```
metrics = history.history
plt.figure(figsize=(16,6))
plt.subplot(1,2,1)
plt.plot(history.epoch, metrics['loss'], metrics['val_loss'])
plt.legend(['loss', 'val_loss'])
plt.ylim([0, max(plt.ylim())])
plt.xlabel('Epoch')
plt.ylabel('Loss [CrossEntropy]')

plt.subplot(1,2,2)
plt.plot(history.epoch, 100*np.array(metrics['accuracy']), 100*np.array(metrics['val_accuracy']))
plt.legend(['accuracy', 'val_accuracy'])
plt.ylim([0, 100])
plt.xlabel('Epoch')
plt.ylabel('Accuracy [%]')
```

```
[25] model.evaluate(test_spectrogram_ds, return_dict=True)

13/13 ————— 2s 158ms/step - accuracy: 0.8272 - loss: 0.5049
{'accuracy': 0.8425480723381042, 'loss': 0.4785589575767517}
```

```
[26] y_pred = model.predict(test_spectrogram_ds)

13/13 ————— 1s 44ms/step
```

```
[27] y_pred = tf.argmax(y_pred, axis=1)
```

```
[28] y_true = tf.concat(list(test_spectrogram_ds.map(lambda s,lab: lab)), axis=0)
```

```
confusion_mtx = tf.math.confusion_matrix(y_true, y_pred)
plt.figure(figsize=(10, 8))
sns.heatmap(confusion_mtx,
            xticklabels=label_names,
            yticklabels=label_names,
            annot=True, fmt='g')
plt.xlabel('Prediction')
plt.ylabel('Label')
plt.show()
```

```
x = data_dir/'no@01bb6a2a_nohash_0.wav'
x = tf.io.read_file(str(x))
x, sample_rate = tf.audio.decode_wav(x, desired_channels=1, desired_samples=16000,)
x = tf.squeeze(x, axis=-1)
waveform = x
x = get_spectrogram(x)
x = x[tf.newaxis,...]

prediction = model(x)
x_labels = ['no', 'yes', 'down', 'go', 'left', 'up', 'right', 'stop']
plt.bar(x_labels, tf.nn.softmax(prediction[0]))
plt.title('No')
plt.show()

display.display(display.Audio(waveform, rate=16000))
```

```
class ExportModel(tf.Module):
    def __init__(self, model):
        self.model = model

    # Accept either a string-filename or a batch of waveforms.
    # You could add additional signatures for a single wave, or a ragged-batch.
    self.__call___.get_concrete_function(
        x=tf.TensorSpec(shape=(), dtype=tf.string))
    self.__call___.get_concrete_function(
        x=tf.TensorSpec(shape=[None, 16000], dtype=tf.float32))

    @tf.function
    def __call__(self, x):
        # If they pass a string, load the file and decode it.
        if x.dtype == tf.string:
            x = tf.io.read_file(x)
            x, _ = tf.audio.decode_wav(x, desired_channels=1, desired_samples=16000,)
            x = tf.squeeze(x, axis=-1)
            x = x[tf.newaxis, :]

        x = get_spectrogram(x)
        result = self.model(x, training=False)

        class_ids = tf.argmax(result, axis=-1)
        class_names = tf.gather(label_names, class_ids)
        return {'predictions': result,
                'class_ids': class_ids,
                'class_names': class_names}
```




```
[32] export = ExportModel(model)
export(tf.constant(str(data_dir/'no/01bb6a2a_nohash_0.wav')))

{'predictions': <tf.Tensor: shape=(1, 8), dtype=float32, numpy=
array([[ 2.4269834,  3.6909025, -2.406927,  4.960789, -2.9688535,
        -2.1688123, -2.7355509, -1.0081533]], dtype=float32)>,
'class_ids': <tf.Tensor: shape=(1,), dtype=int64, numpy=array([3])>,
'class_names': <tf.Tensor: shape=(1,), dtype=string, numpy=array([b'no'], dtype=object)>}
```

```
tf.saved_model.save(export, "saved")
imported = tf.saved_model.load("saved")
imported(waveform[tf.newaxis, :])

{'class_ids': <tf.Tensor: shape=(1,), dtype=int64, numpy=array([3])>,
'predictions': <tf.Tensor: shape=(1, 8), dtype=float32, numpy=
array([[ 2.4269834,  3.6909025, -2.406927,  4.960789, -2.9688535,
        -2.1688123, -2.7355509, -1.0081533]], dtype=float32)>,
'class_names': <tf.Tensor: shape=(1,), dtype=string, numpy=array([b'no'], dtype=object)>}
```

Output:



Conclusion: Hence, the experiment is conducted successfully and the results have been verified.

EXPERIMENT - 11

Aim:

Text generation using RNN.

Experiment Conducted:

```
[ ] import tensorflow as tf

import numpy as np
import os
import time
```

```
[ ] path_to_file = tf.keras.utils.get_file('shakespeare.txt', 'https://storage.googleapis.com/download.tensorflow.org/data/shakespeare.txt')

Downloading data from https://storage.googleapis.com/download.tensorflow.org/data/shakespeare.txt
1115394/1115394 [=====] - 85 0us/step
```

```
[ ] # Read, then decode for py2 compat.
text = open(path_to_file, 'rb').read().decode(encoding='utf-8')
# length of text is the number of characters in it
print(f'Length of text: {len(text)} characters')

Length of text: 1115394 characters
```

```
• # Take a look at the first 250 characters in text
print(text[:250])
```

```
First Citizen:
Before we proceed any further, hear me speak.

All:
Speak, speak.

First Citizen:
You are all resolved rather to die than to famish?

All:
Resolved. resolved.

First Citizen:
First, you know Caius Marcius is chief enemy to the people.
```

```
[ ] # The unique characters in the file
vocab = sorted(set(text))
print(f'{len(vocab)} unique characters')

65 unique characters
```

```
[ ] example_texts = ['abcdefg', 'xyz']

chars = tf.strings.unicode_split(example_texts, input_encoding='UTF-8')
chars

<tf.RaggedTensor [[b'a', b'b', b'c', b'd', b'e', b'f', b'g'], [b'x', b'y', b'z']]>
```

```
[ ] ids_from_chars = tf.keras.layers.StringLookup(
    vocabulary=list(vocab), mask_token=None)
```

```
[ ] ids = ids_from_chars(chars)
ids

<tf.RaggedTensor [[40, 41, 42, 43, 44, 45, 46], [63, 64, 65]]>
```

```
[ ] chars_from_ids = tf.keras.layers.StringLookup(
    vocabulary=ids_from_chars.get_vocabulary(), invert=True, mask_token=None)
```

```
[ ] chars = chars_from_ids(ids)
chars

<tf.RaggedTensor [[b'a', b'b', b'c', b'd', b'e', b'f', b'g'], [b'x', b'y', b'z']]>
```

```
[ ] tf.strings.reduce_join(chars, axis=-1).numpy()

array([b'abcdefg', b'xyz'], dtype=object)
```

```
[ ] def text_from_ids(ids):
    return tf.strings.reduce_join(chars_from_ids(ids), axis=-1)
```

```
[ ] all_ids = ids_from_chars(tf.strings.unicode_split(text, 'UTF-8'))
all_ids

<tf.Tensor: shape=(1115394,), dtype=int64, numpy=array([19, 40, 67, ..., 46, 9, 1])>
```

```
[ ] ids_dataset = tf.data.Dataset.from_tensor_slices(all_ids)
```

```
[ ] for ids in ids_dataset.take(10):
    print(chars_from_ids(ids).numpy().decode('utf-8'))
```

```
F
i
r
s
t

C
i
t
i
```

```
[ ] seq_length = 100
```

```
[ ] sequences = ids_dataset.batch(seq_length=1, drop_remainer=True)

for seq in sequences.take(1):
    print(chars_from_ids(seq))

tf.Tensor(
[b'f' b'i' b'r' b's' b't' b' ' b'c' b'i' b't' b'i' b'z' b'e' b'n' b'i'
 b'ln' b'p' b'e' b'f' b'o' b'r' b'e' b' ' b'w' b'e' b' ' b'p' b'r' b'o'
 b'c' b'e' b'e' b'd' b' ' b'a' b'e' b'y' b' ' b'f' b'u' b'r' b't' b'h'
 b'e' b'r' b'i' b' ' b'h' b'e' b'a' b'r' b' ' b'm' b'e' b' ' b's' b'p'
 b'e' b'a' b'k' b' ' b'ln' b'ln' b'a' b'i' b'i' b'i' b'ln' b's' b'p' b'e'
 b'e' b'k' b' ' b' ' b's' b'p' b'e' b'a' b'k' b' ' b'ln' b'ln' b'r' b'i'
 b'r' b's' b't' b' ' b'c' b'i' b't' b'i' b'z' b'e' b'n' b'i' b'ln' b'y'
 b'o' b'u' b' ' ], shape=(101,1), dtype=string)

[ ] for seq in sequences.take(5):
    print(text_from_ids(seq).numpy())

b'First Citizen:\nBefore we proceed any further, hear me speak.\n\nAll:\nSpeak, speak.\n\nFirst Citizen:\nYou '
bare all resolved rather to die than to finish?\n\nAll:\nResolved, resolved.\n\nFirst Citizen:\nFirst, you k'
b'Now Caius Marcius is chief enemy to the people.\n\nAll:\nWe know't, we know't.\n\nFirst Citizen:\nLet us k'
b'Il him, and we'll have corn at our own price.\n\ns't a verdict?\n\nAll:\nMo more talking on't; let it be d'
b'one: away, away!\n\nSecond Citizen:\nOne word, good citizens.\n\nFirst Citizen:\nWe are accounted poor citi'
```

```
[ ] def split_input_target(sequence):
    input_text = sequence[:-1]
    target_text = sequence[1:]
    return input_text, target_text

[ ] split_input_target(list("TensorFlow"))

([['t', 'e', 'n', 's', 'o', 'r', 'f', 'l', 'o'],
 ['e', 'n', 's', 'o', 'r', 'f', 'l', 'o', 'w']])

[ ] dataset = sequences.map(split_input_target)

[ ] for input_example, target_example in dataset.take(1):
    print("Input :", text_from_ids(input_example).numpy())
    print("Target:", text_from_ids(target_example).numpy())

Input : b'First citizen:\nBefore we proceed any further, hear me speak.\n\nAll:\nSpeak, speak.\n\nFirst Citizen:\nYou '
Target: b'First citizen:\nBefore we proceed any further, hear me speak.\n\nAll:\nSpeak, speak.\n\nFirst Citizen:\nYou '
```

```
[ ] # Batch size
BATCH_SIZE = 64

# Buffer size to shuffle the dataset
# (TF data is designed to work with possibly infinite sequences,
# so it doesn't attempt to shuffle the entire sequence in memory. Instead,
# it maintains a buffer in which it shuffles elements).
BUFFER_SIZE = 10000

dataset = (
    dataset
    .shuffle(BUFFER_SIZE)
    .batch(BATCH_SIZE, drop_remainer=True)
    .prefetch(tf.data.experimental.AUTOTUNE))

dataset

<_PrefetchDataset element_spec=(TensorSpec(shape=(64, 100), dtype=tf.int64, name=None), TensorSpec(shape=(64, 100), dtype=tf.int64, name=None))>
```

```
[ ] # Length of the vocabulary in StringLookup Layer
vocab_size = len(ids_from_chars.get_vocabulary())

# The embedding dimension
embedding_dim = 256

# Number of RNN units
rnn_units = 1024

[ ] class MyModel(tf.keras.Model):
    def __init__(self, vocab_size, embedding_dim, rnn_units):
        super().__init__(self)
        self.embedding = tf.keras.layers.Embedding(vocab_size, embedding_dim)
        self.gru = tf.keras.layers.GRU(rnn_units,
                                         return_sequences=True,
                                         return_state=True)
        self.dense = tf.keras.layers.Dense(vocab_size)

    def call(self, inputs, states=None, return_state=False, training=False):
        x = inputs
        x = self.embedding(x, training=training)
        if states is None:
            states = self.gru.get_initial_state(x)
        x, states = self.gru(x, initial_state=states, training=training)
        x = self.dense(x, training=training)

        if return_state:
            return x, states
        else:
            return x

[ ] model = MyModel(
    vocab_size=vocab_size,
    embedding_dim=embedding_dim,
    rnn_units=rnn_units)

[ ] for input_example_batch, target_example_batch in dataset.take(1):
    example_batch_predictions = model(input_example_batch)
    print(example_batch_predictions.shape, " (batch_size, sequence_length, vocab_size)")

(64, 100, 66) (batch_size, sequence_length, vocab_size)
```

```
[ ] model.summary()

Model: "my_model"
Layer (type)                 Output Shape              Param #
-----
embedding (Embedding)         multiple                   16096
gru (GRU)                     multiple                  3930304
dense (Dense)                 multiple                   67650
-----
Total params: 4022850 (15.35 MB)
Trainable params: 4022850 (15.35 MB)
Non-trainable params: 0 (0.00 B)

[ ] sampled_indices = tf.random.categorical(example_batch_predictions[0], num_samples=1)
sampled_indices = tf.squeeze(sampled_indices, axis=-1).numpy()
```



```
[ ] sampled_indices

array([31,  8, 36, 39, 54, 15, 57, 64, 21, 18, 53, 44, 26,  7, 24, 64, 35,
       14, 56, 25, 10, 25, 41, 10, 27, 46, 49, 40, 32, 10,  8, 19, 22, 13,
       28,  5, 49, 62, 24, 13, 46,  7,  7, 35,  5, 20, 24, 62, 65, 25, 19,
       65, 35, 34, 62,  5, 23, 30, 41, 21, 16, 65, 33, 54, 57, 37, 41, 58,
       47, 21, 54, 16,  9, 61, 14,  0, 43, 58, 51, 33,  2, 44, 42, 26,  2,
       12, 10, 28,  8, 10, 26, 54,  5, 51, 52,  8, 34, 10, 55, 30])

[ ] print("Input:\n", text_from_ids(input_example_batch[0]).numpy())
print()
print("Next Char Predictions:\n", text_from_ids(sampled_indices).numpy())

Input:
b's follows, if you will not change your purpose\nbut undergo this flight, make for Sicilia,\nand there '
Next Char Predictions:
b'R-WZ0BryHEnH,KYVAQL3LB3NgJ53-F1P08jwK7g,,V6KwZLFZVUWL3QMcZT0rxshHOC.vA[UWK]d5IT eCH ;30-3H08lm-U3pQ'
```

```
[ ] loss = tf.losses.sparse_categorical_crossentropy(from_logits=True)

[ ] example_batch_mean_loss = loss(target_example_batch, example_batch_predictions)
print("Prediction shape: ", example_batch_predictions.shape, " # (batch_size, sequence_length, vocabab_size)")
print("Mean loss: ", example_batch_mean_loss)

Prediction shape: (64, 100, 66) # (batch_size, sequence_length, vocabab_size)
Mean loss: tf.Tensor(4.189323, shape=(), dtype=float32)
```

```
[ ] tf.exp(example_batch_mean_loss).numpy()

65.9781
```

```
[ ] model.compile(optimizer='adam', loss=loss)
```

```
[ ] # Directory where the checkpoints will be saved
checkpoint_dir = './training_checkpoints'
# Name of the checkpoint files
checkpoint_prefix = os.path.join(checkpoint_dir, "ckpt_{epoch}")

checkpoint_callback = tf.keras.callbacks.ModelCheckpoint(
    filepath=checkpoint_prefix,
    save_weights_only=True)
```

```
① EPOCHS = 20
```

```
① history = model.fit(dataset, epochs=EPOCHS, callbacks=[checkpoint_callback])
```

```
② Epoch 1/20
172/172 [=====] - 13s 51ms/step - loss: 2.7300
```

```
① class OneStep(tf.keras.Model):
    def __init__(self, model, chars_from_ids, ids_from_chars, temperature=1.0):
        super().__init__()
        self.temperature = temperature
        self.model = model
        self.chars_from_ids = chars_from_ids
        self.ids_from_chars = ids_from_chars

        # Create a mask to prevent "[UNK]" from being generated.
        skip_ids = self.ids_from_chars([ "[UNK]" ])[0], None]
        sparse_mask = tf.sparse_tensor(
            # Put a -inf at each bad index.
            values=-1e9 * tf.ones([len(skip_ids)] * len(skip_ids)),
            indices=skip_ids,
            # Match the shape to the vocabulary
            dense_shape=[len(ids_from_chars.get_vocabulary())])
        self.prediction_mask = tf.sparse.to_dense(sparse_mask)

    @tf.function
    def generate_one_step(self, inputs, states=None):
        # Convert strings to token IDs.
        input_chars = tf.strings.split(inputs, 'UTF-8')
        input_ids = self.ids_from_chars(input_chars).to_tensor()

        # Run the model.
        # predicted_logits.shape is [batch, char, next_char_logits]
        predicted_logits, states = self.model(input_ids=input_ids, states=states,
                                              return_state=True)

        # Only use the last prediction.
        predicted_logits = predicted_logits[:, -1, :]
        predicted_logits = predicted_logits / self.temperature
        # Apply the prediction mask: prevent "[UNK]" from being generated.
        predicted_logits = predicted_logits * self.prediction_mask

        # Sample the output logits to generate token IDs.
        predicted_ids = tf.random.categorical(predicted_logits, num_samples=1)
        predicted_ids = tf.squeeze(predicted_ids, axis=-1)

        # Convert from token IDs to characters
        predicted_chars = self.chars_from_ids(predicted_ids)

        # Return the characters and model state.
        return predicted_chars, states
```

```
[ ] one_step_model = OneStep(model, chars_from_ids, ids_from_chars)
```

```
① start = time.time()
states = None
next_char = tf.constant(['ROME0:'])
result = [next_char]

for n in range(1000):
    next_char, states = one_step_model.generate_one_step(next_char, states=states)
    result.append(next_char)

result = tf.strings.join(result)
end = time.time()
print(result[0].numpy().decode('utf-8'), '\n\n' + '_' * 80)
print('\nRun time:', end - start)
```

```
[ ] start = time.time()
states = None
next_char = tf.constant(['ROME0:', 'ROME0:', 'ROME0:', 'ROME0:', 'ROME0:'])
result = [next_char]

for n in range(1000):
    next_char, states = one_step_model.generate_one_step(next_char, states=states)
    result.append(next_char)

result = tf.strings.join(result)
end = time.time()
print(result, '\n\n' + '_' * 80)
print('\nRun time:', end - start)
```

```
[ ] tf.saved_model.save(one_step_model, 'one_step')
one_step_reloaded = tf.saved_model.load('one_step')

WARNING:tensorflow:Skipping full serialization of Keras layer <_main__OneStep object at 0x7c263b9a5580>, because it is not built.
WARNING:tensorflow:Model's '__init__()' arguments contain non-serializable objects. Please implement a 'get_config()' method in the subclassed Model for proper saving and loading. Defaulting to empty config.
WARNING:tensorflow:Model's '__init__()' arguments contain non-serializable objects. Please implement a 'get_config()' method in the subclassed Model for proper saving and loading. Defaulting to empty config.

[ ] states = None
next_char = tf.constant(['ROMEO:'])
result = [next_char]

for n in range(100):
    next_char, states = one_step_reloaded.generate_one_step(next_char, states=states)
    result.append(next_char)

print(tf.strings.join(result)[0].numpy().decode("utf-8"))

[ ] class CustomTraining(MyModel):
    @tf.function
    def train_step(self, inputs):
        inputs, labels = inputs
        with tf.GradientTape() as tape:
            predictions = self(inputs, training=True)
            loss = self.loss(labels, predictions)
            grads = tape.gradient(loss, model.trainable_variables)
            self.optimizer.apply_gradients(zip(grads, model.trainable_variables))

        return {'loss': loss}

[ ] model = CustomTraining(
    vocab_size=len(ids_from_chars.get_vocabulary()),
    embedding_dim=embedding_dim,
    rnn_units=rnn_units)

[ ] model.compile(optimizer = tf.keras.optimizers.Adam(),
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True))

[ ] model.fit(dataset, epochs=1)

172/172 [=====] - 13s 52ms/step - loss: 2.7221
<keras.src.callbacks.History at 0x7c2611b79660>
```

```
Ⓜ EPOCHS = 10

mean = tf.metrics.Mean()

for epoch in range(EPOCHS):
    start = time.time()

    mean.reset_states()
    for (batch_n, (inp, target)) in enumerate(dataset):
        logs = model.train_step([inp, target])
        mean.update_state(logs['loss'])

        if batch_n % 50 == 0:
            template = f"Epoch {epoch+1} Batch {batch_n} Loss {logs['loss']:.4f}"
            print(template)

    # saving (checkpoint) the model every 5 epochs
    if (epoch + 1) % 5 == 0:
        model.save_weights(checkpoint_prefix.format(epoch=epoch))

    print()
    print(f'Epoch {epoch+1} Loss: {mean.result().numpy():.4f}')
    print(f'Time taken for 1 epoch {time.time() - start:.2f} sec')
    print("=="*80)

    model.save_weights(checkpoint_prefix.format(epoch=epoch))
```

Output:

```
Ⓜ ROMEO:
Two lords that bears that title substancially
Hath made her. Come, lend me with your respects,
And we are but one indeed, and stuff'd
With poltrovs warm.

PAULINA:
I'll pass here;
How many of she, my prophecies, answer,
To bake awak'd to some unwillings.

AUFIDIUS:
There's a woman on the earth.

NORTHUMBERLAND:
Save your own bond fetch me her time?

SEBASTIAN:
A daughter of mine he rounds not to be found,
I mean of moubting hatren hast togster.

MERCUTIONA:
Your brother's son shall be to aspire them
To good his hounds and had shrone-thing.

DUKE OF YORK:
Vixe unsweary home and happoser writ you.

DUKE VINCENTIO:
There's please you bring me to the bourney husband
And made the house of soldiers. Hastings, spent!
As thou mayst tell him Aury, stand by ambments.

ELBOW:
Is the gods be the bloody king!

KATHARINA:
Their mother? why, how now! what's the matter, sir;
But for this night, which going
A beauty of a bleeding smade.
Go with my heart. My very wisdom to
Then break the flinty times to
```

Run time: 3.103527784347534

Run time: 3.146984338760376

Epoch 7 Loss: 1.2877
Time taken for 1 epoch 10.32 sec

Epoch 8 Loss: 1.2460
Time taken for 1 epoch 10.28 sec

Epoch 9 Loss: 1.2067
Time taken for 1 epoch 10.30 sec

Epoch 10 Loss: 1.1672
Time taken for 1 epoch 10.51 sec

Conclusion: Hence, the experiment is conducted successfully and the results have been verified.